
Semantic Foundations for the Analysis of Program Revisions

Dakota Bryan

Project Advised by Evan Chang and
Matthew Hammer

Program Analysis: Theory and Practice,
Fall 2025

Motivation

Developers have informal notions of the semantics of a program revision. We aim to formalize these semantics and provide an abstract interpretation framework to reason about them.

Understanding the Problem

How can we **describe** the semantic effect of a program revision? (Concrete / collecting semantics)

How can we **reason** about these effects? (Abstract semantics)

Contributions and Outline

Part 1: We **describe** the semantic effect of a program revision by providing concrete / collecting semantics

Part 2: We give the requirements for **reasoning** about these effects by providing parameterized abstract semantics

Semantics of Program Revisions

How can we describe the semantic effect of a program revision?

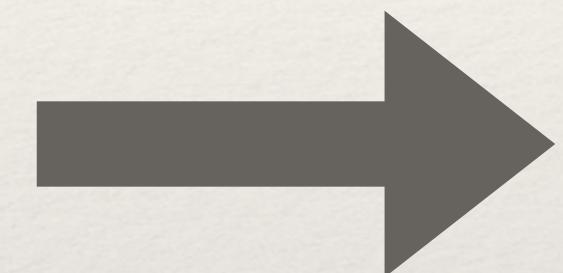
How can we reason about these effects?

Idea: Reachable states of the “old” program are either **removed** or **related**.
Reachable states of the “new” program are either **added** or **related**.

Remaining Question: When is an “old” / “new” program state removed / added?

Example

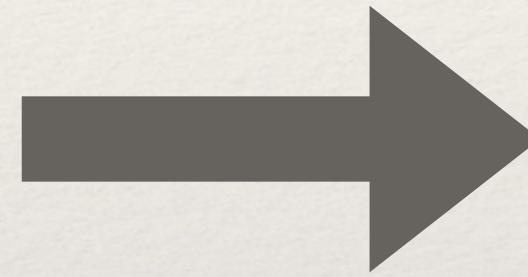
```
x := nonDet()
if (x > 0)
    sgn := 1
else
    sgn := -1
```



```
x := nonDet()
if (x > 0)
    sgn := 1
else
    sgn := -1
if (x = 0)
    sgn := 0
```

Example: Entire State

```
x := nonDet()  
if (x > 0)  
    sgn := 1  
else  
    sgn := -1
```



```
x := nonDet()  
if (x > 0)  
    sgn := 1  
else  
    sgn := -1  
if (x = 0)  
    sgn := 0
```

$x \mapsto 0, \text{sgn} \mapsto -1$

$x \mapsto 0, \text{sgn} \mapsto 0$

Example: State Projection (sgn)

```
x := nonDet()  
if (x > 0)  
    sgn := 1  
else  
    sgn := -1
```



```
x := nonDet()  
if (x > 0)  
    sgn := 1  
else  
    sgn := -1  
if (x = 0)  
    sgn := 0
```

$x \mapsto 0, \text{sgn} \mapsto 0$

Solution: State Correspondence

- ❖ A “new” state, σ_2 , is added if for any reachable “old” state, σ_1 , $\sigma_2 \nsim \sigma_1$
- ❖ \sim is a user defined relation between “old” states and “new” states

state correspondence	$s ::= p_1, p_2$
state projection	$p ::= \circ \mid p, x$
program variables	x

$$s = (\text{sgn}), (\text{sgn})$$

$$\sigma_1 \sim \sigma_2 \iff \sigma_1(\text{sgn}) = \sigma_2(\text{sgn})$$

Solution: State Correspondence

How can the effect of a program revision be described? (Concrete / collecting semantics)

How can we reason about these effects?

Solution: Reachable states of the “old” and “new” programs are characterized as **removed**, **related**, or **added** (concretely, sets) via a state correspondence

Abstraction

How can the effect of a program revision be described?

How can we reason about these effects?

Idea: Over-approximate **removed**, **related**, and added sets of states

Intricacies of Abstraction: Partitioning

Recall: $\text{removed} = \{\sigma_1 \in \Sigma_1 \mid \forall \sigma_2 \in \Sigma_2 . \sigma_1 \not\sim \sigma_2\}$

Consider: $\hat{\sim}$ and $\hat{\sigma}$ where $\gamma(\hat{\sigma}) = \Sigma$

Challenge: If $\hat{\sigma}$ merges all abstract states into one element, then either all states are removed or none of them are

Intricacies of Abstraction: Partitioning

One Solution: $\hat{\sigma}$ is split into elements representing subsets of $\gamma(\hat{\sigma})$

$$\text{removed} = \{\hat{\sigma}_i \in \hat{\sigma}_1 \mid \forall \hat{\sigma}_j \in \hat{\sigma}_2 . \hat{\sigma}_i \hat{\not\sim} \hat{\sigma}_j\}$$

Intricacies of Abstraction: Under-Approximation

$$\text{removed} = \{\hat{\sigma}_i \in \hat{\sigma}_1 \mid \forall \hat{\sigma}_j \in \hat{\sigma}_2 . \hat{\sigma}_i \hat{\approx} \hat{\sigma}_j\}$$

Problem: If $\Sigma_2 \subseteq \gamma(\hat{\sigma}_2)$, then there are “more opportunities” for $\hat{\sigma}_i \hat{\approx} \hat{\sigma}_j$ for some $\hat{\sigma}_i$ resulting in the possibility for **removed** to be an **under-approximation**

Intricacies of Abstraction: Under-Approximation

$$\text{removed} = \{\hat{\sigma}_i \in \hat{\sigma}_1 \mid \forall \hat{\sigma}_j \in \hat{\sigma}_2 . \hat{\sigma}_i \not\approx \hat{\sigma}_j\}$$

Solution: We require that the abstraction of reachable stats for the “new” program is an **under-approximation** ($\gamma(\hat{\sigma}_2) \subseteq \Sigma_2$)

Parameterized Abstract Semantics

How can the effect of a program revision be described? (Concrete / collecting semantics)

How can we reason about these effects? (Parameterized abstract semantics)

We require:

- ❖ An over- and under-approximating abstraction of reachable states,
- ❖ In a finitely partitioned abstract domain,
- ❖ With the state correspondence operation lifted to this domain

Thank you, Questions?